

Chuzhditsa: a featural loanword register for Cyrillic, from Bulgarian to the Slavic superset

Author name withheld for review*

Abstract

Japanese distinguishes loanwords at a glance: katakana is a parallel register whose visual form alone announces foreignness, and whose adaptation conventions, while not strictly deterministic, are stable enough that a specified source pronunciation yields a largely predictable spelling. No major Slavic Cyrillic orthography has a visually distinct, productive register dedicated to foreign phonology as such, and the result is systematic indeterminacy (Bulgarian **ВОШИНГТОН** ~ **УОШИНГЪН**, Данчев 1982; Serbian and Russian counterparts abound). We present *Chuzhditsa* (Bulg. чуждица, ‘foreignism’), a computable, Cyrillic-based *phonological citation register* for foreign words, names, and pedagogical transcription; its secondary use is as a visual register for unassimilated loans. The system is constructed from attested Cyrillic material and evaluated for internal consistency, cross-linguistic coverage, and typographic implementability. It supplies (i) a featural grammar of diacritic operations deriving letters for non-Slavic phonemes, (ii) a prosodic layer for stress, tone, and length, and (iii) a register-marking typeface whose Glagolitic-derived forms perform the katakana function visually. The letter inventory is assembled conservatively: revivals from historical Bulgarian orthography (the yuses **Ѧ ѧ**) alongside the living Macedonian **Ѕ**, imports from the wider Cyrillic ecosystem (Kazakh **Ғ Ң**, Bashkir **Җ**, Tajik **Ӣ**), and rule-derived residue. We evaluate the system against the phonologies of twenty major languages, each under a declared reference variety and a three-way coverage classification (exact, conventionally merged, declared loss), generalize it from its Bulgarian pilot to the full Slavic Cyrillic superset (with deterministic bridges from Latin-alphabet Slavic languages and an external validation on Albanian), and describe a parametric implementation as an OpenType typeface family in which the fusion and anchoring rules of the orthography are executable GSUB/GPOS code. We argue that the laryngeal series (**Ѧ ѧ Ѩ** plus the aspiration modifier **Ѧ̣**) illustrates a general principle for extending alphabets: articulatory gradients should map to visual gradients.

1 Introduction

Every literate community borrows words faster than its orthography naturalizes them — an instance of the general problem of fit between a writing system and a changing language (Coulmas 2003). Japanese is unusual in having institutionalized the problem: katakana is a full parallel syllabary whose use *is* the marking of foreignness, and whose adaptation conventions ([nart] → ナイト) are conventionalized enough that, for a given source pronunciation, writers usually converge on one spelling. No major Slavic Cyrillic orthography has an analogous visual register. Loanwords enter Bulgarian, Russian, Serbian or Ukrainian through ad hoc transcription: sound-based and spelling-based renderings compete (**ЛАНЧ** ~ **ЛЪНЧ** for *lunch*), phonemes without Cyrillic letters collapse unpredictably (/θ/ → Т or С; /w/ → В or У), and nothing in the written form signals that a word is a guest rather than a native.

This paper describes Chuzhditsa, a designed orthographic register that supplies the missing device for Cyrillic, piloted on Bulgarian and then generalized to the Slavic superset.

One distinction must be drawn before anything else, because three different tasks hide under “writing foreign words,” and a system is judged unfairly when measured against a task it does not attempt.

*Fonts, build toolchain, and the twenty-language audit are available in the accompanying repository; the typeface used for all object-language examples in this manuscript is the system being described.

Loanword adaptation bends a word to the recipient phonology — katakana’s main business, and the business of every integrated borrowing. *Phonemic transcription* records the source phonology in a technical alphabet — the IPA’s business. Between them sits *phonological citation*: writing a foreign word, name, or example so that a native reader can recover an adequate approximation of its source pronunciation, in the native script, without technical training. Chuzhditsa is primarily the third: unlike katakana it *extends* the recipient inventory to preserve source contrasts rather than collapsing them into native phonotactics; unlike the IPA it stays inside Cyrillic reading habits. Its use as a marking register for not-yet-naturalized loans is a secondary, optional application, and a loan that naturalizes is expected to leave the register. The system makes three claims of general interest to writing-system design. First, that a loanword register can be *featural* in Sampson’s (1985) sense: rather than inventing letters ad libitum, a small set of diacritic operations — each meaning exactly one phonetic feature — derives the entire extension inventory, echoing the compositional logic of Hangul (Sampson 1985; Kim-Renaud 1997). Second, that such a register can be assembled almost entirely from attested material: dead letters of the target orthography’s own history and living letters of related orthographies, with invention reserved for genuine residue. Third, that the register-marking function that Japanese distributes over a second character inventory can instead be carried by a second *typeface register* — realized here as an OpenType family whose substitution (GSUB) and positioning (GPOS) rules are the orthography’s fusion and attachment rules, executable rather than merely described.

2 Background

2.1 Katakana as an adaptation layer

Katakana’s success rests on three properties (Irwin 2011; Stanlaw 2004): conventionalized conversion that renders the loan’s *adapted Japanese pronunciation* — an adaptation that has historically travelled by ear, by dictionary, and sometimes by spelling (Irwin 2011: 76–79) — rather than transliterating the source orthography letter for letter, the orthographic reflex of loanword adaptation as a systematic mapping onto native categories, whether modelled phonologically or perceptually (Kang 2011; Peperkamp & Dupoux 2003); graphic distinctness sufficient to mark register at a glance; and extensibility (ヴ for /v/, ファ for /fa/) when the inventory runs out. Chuzhditsa adopts all three as requirements, and hardens the first: its input is pronunciation only. The idealization should be stated as such — real katakana practice varies by borrowing period, donor language, brand convention, and register (Irwin 2011) — and Chuzhditsa aims to be what katakana is often described as being: the target is a convention strict enough to converge, not a claim that Japanese practice already converges. Katakana also carries a social record that a designed analogue cannot disclaim: loanword marking has been read as a badge of incomplete nativeness, and katakana is deployed to mark foreigners’ accented Japanese as well as foreign words (Irwin 2011; Robertson 2020). We therefore state the register’s semantics narrowly: it asserts that a word’s phonology lies off the native grid — nothing about its right to exist in the language, and nothing about its user. Its proposed semantics is lexical rather than ethnolinguistic — it marks words, never speakers — though a designed intention is not a sociolinguistic guarantee: registers acquire social meanings their designers do not control, and whether readers would honor the boundary is an empirical question for user studies. A naturalized loan is expected to exit the register, exactly as テンプラ became 天ぷら.

2.2 The Cyrillic commonwealth

Cyrillic today serves more than fifty languages, whose principal non-Slavic adaptations Comrie (1996) surveys, and the non-Slavic ones have already engineered letters for most sounds Slavic lacks: Kazakh қ /q/, ң /ŋ/, Ғ /ɣ/; Bashkir ҫ /θ/, ҫ /ð/; Tajik Ӣ /i:/ and Ҳ (/h/ in Tajik, reassigned here to /ħ/; the borrowing is of the *grapheme and its featural construction* — x plus retraction — not of its Tajik value, and the Caucasian digraph xI is the nearer phonetic precedent); Altai Ӗ /ø/, Ӛ /y/; Belarusian ӱ /w/; Serbian ӱ /dʒ/;

the Caucasian palochka I. These letters carry decades of typographic and pedagogic practice. A design that imports them inherits that practice for free — an option unavailable to katakana’s designers and, we will argue, the single largest economy in the present system.

2.3 The Glagolitic precedent

Bulgaria is a primary locus of a script replacement: Glagolitic, the first Slavic alphabet, created — or, on Cubberley’s reading, formalized — by Constantine-Cyril for the Moravian mission of 863, was gradually displaced by Cyrillic, devised by Constantine’s disciples at the Preslav Literary School in the 890s on the model of Greek uncial, drawing on Glagolitic mainly for the sounds Greek lacked (Schenker 1995; Cubberley 1996). Glagolitic then survived for a millennium as a special-register script (Croatian liturgical use into the twentieth century; the last missal printed in Glagolitic script appeared in 1905, later service books shifting to Latin transliteration). Chuzhditsa knowingly reverses the vector: the visual DNA of the *oldest* Slavic letterforms — round bowls, unicameral proportions — marks the *newest* words, and the letters it revives (А А А А) inherit the sound-slots of Glagolitic ancestors (the Cyrillic shapes themselves are reworkings).

3 Design principles

1. **Determinism.** One *specified input* phoneme, one rendering, within a declared mapping: given the same source pronunciation in the reference variety, two transcribers converge. The choice of source pronunciation itself (dialect, variant, spelling pronunciation) is outside the system’s jurisdiction, and loss is permitted but must be declared (§9).
2. **Featural transparency.** Every diacritic operation encodes one recurrent phonological *relation* to its base letter and applies wherever that relation appears. Some operations are featurally atomic (nasalization, aspiration, length); one is relational rather than atomic — the retraction hook means “place shifted off the native grid,” interpreted relative to the base (§4) — so the system is featural in Sampson’s extended sense rather than the strict one.
3. **Pair symmetry.** Obstruent letters enter the system in voiceless/voiced pairs aligned with the native series. This principle outranks heritage: the historical theta-letter fita (Θ) was rejected because it has no voiced partner and derives from no rule; the Bashkir pair җ/ӓ (/θ//ð/) does its work systematically.
4. **Precedent first.** Revive native dead letters, then import living Cyrillic letters, and only then derive new combinations; never invent an unattested base shape.
5. **Pan-Cyrillic legibility.** Every base skeleton follows the reader’s existing Cyrillic, and every extension must be guessable at first sight by any Cyrillic reader from its base letter plus a transparent mark — a design criterion the orthography-development literature calls ease of transfer (Cahill & Rice 2014: 9–25).
6. **Register by typeface.** Foreignness is marked by letterform register, not by punctuation or a second inventory: text set in the Chuzhditsa face is thereby marked, exactly as katakana text is — in the terms of the biscriptality literature, a *diaphasic diglyphia*: functional differentiation between two glyphic variants of one script (Bunčić, Lippert & Rabus 2016; the nearest historical parallel is the German practice of setting foreign words in roman within blackletter text), and a socially meaningful orthographic choice in Sebba’s (2007) sense.

4 The featural grammar

Table 1 gives the full operation set. The retraction hook generalizes a precedent internal to the Slavic base inventory itself — **ѡ** is **ѡ** with a hook — and the glide breve generalizes **и** → **й**. Where Unicode supplies precomposed hooked letters (**ѣ ѥ ѧ Ѩ ѩ**) they are used; elsewhere the same hook is written with combining U+0322 (**Ѣ ѣ Ѥ ѥ Ѧ ѧ Ѩ ѩ**).

Mark	Operation	Derivations
◌̆ breve	vowel → glide	и → й /j/ · у → ѳ /w/
◌̇ hook below	place shifted off the native grid	ѣ /q/ · ѥ /h/ · ѧ /ŋ/ · Ѩ /θ/ · ѩ /ð/ · Ѫ /t/ · ѫ /d/ · Ѭ /tʃ/
◌̈ crossbar	lenition	г → ѣ /x/
◌̊ double dot	vowel fronting	Ѧ /æ/ · ѧ /ø/ · Ѩ /y/ · ѩ /e/
◌̋ yus tail	nasalization	Ѧ /ɔ̃/ · ѧ /ɛ̃/ · Ѩ /ɑ̃/ · ѩ /œ̃/
◌̌ ◌̍ raised h	aspiration; breathy voice	п̌ т̌ к̌ ч̌ · б̌ а̌ г̌ у̌
◌̄ macron	vowel length	Ѣ ѣ (precomposed) · Ѥ ѥ Ѧ
І palochka	laryngeal–pharyngeal action	Ѣ ѣ (ejectives) · bare І = /ʕ/
◌̎	palatalization	Ѣ /j/ · ѧ /ɟ/ · Ѩ /tʃ/ · ѩ /dʒ/

Table 1: The featural operations. Each mark encodes a stable operation; most are featurally atomic, while the hook is relational to its base (§4). Precomposed Unicode letters are used where they exist.

Table 2 lays out the complete designed inventory — every letter the system adds to a Slavic Cyrillic base, set in the typeface itself.

4.1 The laryngeal gradient

The three back fricatives that plague Cyrillic transcription of Arabic, English and German — /x/, /ħ/, /h/ — receive a designed gradient: articulatory weakening maps to graphic lightening. Native **х** (/x/~/ħ/) is two full crossing strokes, the densest friction; pharyngeal **х̌** (/ħ/) is **х** pulled deeper by the retraction hook; glottal **х̍** (/h/~/ɦ/) is a single open arc, literally the most open glyph of the three; and aspiration **х̎**, a feature rather than a segment, is **х** reduced to a superscript. Arabic exhibits the full triple within ordinary vocabulary: **х̍** **х̌** **х̎** /hila:l/, **х̍** **х̌** **х̎** /muħammad/, **х̍** **х̌** **х̎** /xa:lɪd/. We suggest the mapping *articulatory gradient* → *visual gradient* as a reusable principle for alphabet extension.

4.2 The Indic stress test

Hindi–Urdu and Bengali cross a four-way laryngeal contrast with a dental/retroflex place contrast. The grammar covers the entire grid with zero new letters: **т** /t/, **т̌** /tʰ/, **ѧ** /d/, **ѧ̌** /dʰ/; retroflex **Ѣ** **ѣ** **Ѧ** **ѧ**, plus **ѣ** /ɳ/, **ѥ** /ɻ/, **ѧ̌** /ɻ̌/. A single word such as **Ѣ̌** **ѣ̌** **ѧ̌** (*thīk*, /tʰi:k/) composes three features on one base — hook, aspiration, length — which is precisely the featural claim made concrete.

5 The ligature stratum

Cyrillic has always ligated: **ѡ** is a historic ligature of **и**+**ѡ** (the **ѡ** element later dropped), Serbian **ѡ** **ѡ** were fused from **л**/**н**+**ѡ** by Vuk Karadžić in 1818, building on Mrkalj’s digraphs, and Old Bulgarian possessed iotated nasals. Chuzhditsa’s ligature stratum is therefore conservative, and it is implemented as executable GSUB rules rather than prose: (i) **ѡ** **ѡ** fuse to **ѡ̌**/**ѡ̌**-forms automatically (**ѡ̌** **ѡ̌**, **ѡ̌** **ѡ̌** set with fused palatals); (ii) the iotated yuses **ѡ̌** /ǰ/ and **ѡ̌** /ǰ/ are revived as directly typeable letters, and the sequences **ѡ̌**+**ѡ̌** and **ѡ̌**+**ѡ̌** fuse to them automatically (French *bien* → **ѡ̌**; Polish *pięć* → **ѡ̌**); (iii) ejective digraphs **Ѣ** **ѣ** **Ѧ** **ѧ** fuse with the palochka *raised to the upper half*, echoing the apostrophe of scholarly romanization ([tʰ kʰ]) — a fused full-height palochka was rejected in testing because **Ѣ**

Letter	Unicode	Value	Construction	Origin
Ÿŷ	U+040E/045E	/w/	y + breve	import: Belarusian
Џџ	U+040F/045F	Ṫḑz/	own shape	import: Serbian
Ҙҙ	U+04AA/04AB	/θ/	c + hook	import: Bashkir
Зз	U+0498/0499	/ð/	z + hook	import: Bashkir
Ққ	U+049A/049B	/q ɢ/	к + hook	import: Kazakh
Ңң	U+04A2/04A3	/ŋ/	н + hook	import: Kazakh
Ҳҳ	U+04B2/04B3	/ħ/	x + hook	import: Tajik
Ғғ	U+0492/0493	/ɣ/	г + crossbar	import: Kazakh
Һһ	U+04BA/04BB	/h ɦ/	own shape	import: Tatar
Ѕѕ	U+0405/0455	Ṫḑz/	own shape	living Macedonian (dzelo)
І	U+04C0	/ʃ/; Ҁ ҁ ҂	palochka	import: Caucasus
Ҁ ҁ ҂	base + U+0322	/t̪ d̪ ɳ̪/	base + hook	derived by rule
ḥ	U+02B0	aspiration	raised miniature Һ	derived
ḥ	U+02B1	breathy voice	raised hooked Һ	derived
Ăă	U+04D2/04D3	/æ/	a + double dot	import: Mari
Ӗӗ	U+04E6/04E7	/ø œ/	o + double dot	import: Altai
Ӣӣ	U+04F0/04F1	/y ʏ/	y + double dot	import: Altai
Ыы	U+042B/044B	/i w/	own shape	import: Russian; also Old Bulgarian
Ӑӑ	U+046A/046B	/õ ö/	nasal o	revival: in use until 1945
Ӓӓ	U+0466/0467	/ē/	nasal e	revival: Old Bulgarian
Ӕӕ	U+0468/0469	/jē/	ligature Ъ+Ӑ	revival: Old Bulgarian
Ӧӧ	U+046C/046D	/jō/	ligature Ъ+Ӓ	revival: Old Bulgarian
Ṫ Ṫ̃	V + U+0328	/ã œ̃/	vowel + yus tail	derived by rule
ӢӢ̃ Ӣ̃Ӣ̃	U+04E2/3, EE/EF	/i: u:/	vowel + macron	import: Tajik
̄	U+0304	length	macron	derived
◌́ ◌̀ ◌̂	U+0301/0300/030C	tone, stress	dictionary layer	derived

Table 2: The complete designed inventory. The letter column is set in the Chuzhditsa face. Every letter is a standard Cyrillic-block codepoint; the combining and modifier characters are likewise standard Unicode (The Unicode Consortium 2024).

became confusable with Һ. The governing rule: only sequences denoting a single phoneme, with Slavic or Old Cyrillic precedent, may fuse; mechanical soft pairs (Ҙ Ҙ̃) remain visibly two-part.

6 The prosodic stratum

Stress takes the grave of Bulgarian lexicography (бѣбѣуѣкѣ); Mandarin tones 2–4 take acute, caron, grave, with tone 1 unmarked — pinyin minus the macron, which Chuzhditsa reserves for length. The unmarked first tone resolves the collision by construction: 𐍌̄ can only mean a long vowel. Three costs are declared rather than hidden. Tone and stress marks are homoglyphic across source languages, resolved only by knowing the lexical source (katakana shares exactly this property). Diacritic stacking degrades: Vietnamese *phở* wants quality, length and tone at once (𐍌̄̂̃), and the running-text rule drops length first (𐍌̂̃). And four marks cannot represent six-tone systems without folding (Vietnamese *hỏi/ngã* merge). The design verdict generalizes katakana’s own: prosody is optional annotation — a dictionary layer, not an orthographic one — and the system claims no tonal adequacy beyond that annotative role.

7 From Bulgarian to the Slavic superset

The pilot treated Bulgarian’s thirty letters as the base. Nothing in the featural grammar is Bulgarian-specific, however, and the generalization is administrative rather than structural: the base becomes

each language’s own Cyrillic alphabet (for the phoneme inventories, see Comrie & Corbett 1993), and the extension and prosodic strata apply unchanged. Concretely, the typeface’s character set was widened to the Slavic Cyrillic superset — Russian *ѐ ё ъ*, Ukrainian *і ї є ґ*, Belarusian *ѣ*, Serbian *ј љ њ ѝ њ џ ѡ Ѣ*, Macedonian *ѕ ґ ґ ґ* — so that any Slavic Cyrillic text sets natively. One design conflict surfaced and is instructive: the pilot’s round e-form was, in fact, the shape of Ukrainian *є*; pan-Slavic coverage forced e to an angular form so that the e/е contrast survives. Extending a script for one language can silently occupy a sibling’s letterform space; superset-first design catches this early.

For each language the system’s marginal contribution differs, and several languages turn out to need the extension layer for *native* material, not only loans: Slovak *ä* is *А* and Slovak *h* is /h/ → *н*; Slovenian’s native schwa (*polglasnik*) is exactly *ь* (*пѣс* /pəs/ ‘dog’); Serbian’s pitch accents are the prosodic layer’s use case; Polish *ą ę* are the yuses meeting their own descendants (*макѡ, пѡчѣ*). Latin-alphabet Slavic languages connect through a deterministic bridge: Gaj’s digraphs map one-to-one onto Serbian Cyrillic letters (*lj nj dž* → *љ њ џ*), and the residue is featural — Czech *ř* /r/, the notorious singleton, is written *ř* with the caron it already wears; Slovak long syllabic liquids *ĺ ř* take the length macron (*вѣба* /vr:ba/).

The generalization also has ancestors to acknowledge: the determinism principle is Vuk Karadžić’s *нѣмѣ као што говориш* avant la lettre, and Serbian orthography contributed *ѣ љ њ ј* — and the very idea of fused palatals — before this project existed; Ukrainian orthography constrained the design as a peer, not a client (the *є* letterform conflict above was resolved in Ukrainian’s favor). The Bulgarian name records the pilot, not a hierarchy.

The citation-register framing predicts something checkable across the family: *integrated* loans differ by language, because each language adapted the word to its own habits, while the *citation* form is shared, because its input is the source pronunciation. Table 3 shows the pattern.

Source	Bulgarian	Russian	Ukrainian	Serbian	shared citation form
<i>weekend</i>	УИКЕНА	УИК-ЭНА	ВІКЕНА	ВИКЕНА	ЎЙКЕНА
<i>jazz</i>	ДЖАЗ	ДЖАЗ	ДЖАЗ	ЏЕЗ	ЏАЗ

Table 3: Integrated loans vary by recipient language; the citation register does not, because its input is the source pronunciation rather than any recipient’s adaptation.

As an external stress test we applied the system to Albanian, a non-Slavic Balkan language: every Albanian phoneme lands on an existing letter (*ë* /ə/ → *ь*, *th/dh* → *џ*, *y* → *ѣ*, *xh* → *ѡ*, *x* /dz/ → *ѕ*, *q/gj* → *кь/гь*): *Шкѣпѣриѡ, Тирѡнѣ*. A system built for Slavic phonology plus the featural operations appears to cover the Balkan sprachbund without amendment. The practical beneficiary is not Albanian itself, which is securely Latin-written, but Cyrillic-language text citing Albanian: the *ë* that Macedonian and Serbian renderings currently drop is exactly *ь*.

8 Typeface implementation

The Chuzhditsa family (Regular, Bold, Italic, Bold Italic; TTF and CFF/OTF) is generated by a ~300-line parametric builder: every glyph is a set of monoline strokes — lines, arcs, and the full circles that give the face its Glagolitic character — expanded to outlines by a union of per-segment rectangles and per-vertex discs, with weight and slant as parameters. Capitals stand at 700 units, lowercase at an x-height of 500 with true two-case architecture (bowl-and-stem *а*, ascending *б*, descending *р*). Every combining mark receives computed GPOS mark-to-base anchors (the macron of *ѡ* attaches 251 units lower than that of *Ѣ*; the retroflex hook of *ѣ* lands on the letter’s stem), and the ligature stratum is a GSUB *liga* feature — the orthography’s rules are the font’s code, verified by HarfBuzz shaping tests in the build toolchain — a concrete instance of Sproat’s (2000) regularity hypothesis: the mapping from linguistic representation to written form is a regular relation, computable by finite-state means. The linguistic point of this engineering is falsifiability: an orthographic rule that exists as a GSUB

lookup either fires or does not, so the specification’s claims about fusion and attachment are machine-checkable properties of an artifact rather than intentions in prose. The manuscript the reader is holding sets all object-language material in this typeface.

9 Evaluation: the twenty-language audit

Method. The audit set is not a demographic top-twenty: it is a high-speaker-count convenience sample adjusted for phonological diversity and Slavic relevance, comprising English, Mandarin, Hindi–Urdu, Spanish, French, Arabic (MSA), Bengali, Portuguese, Russian, Japanese, German, Korean, Vietnamese, Turkish, Italian, Persian, Swahili, Greek, Polish, and Indonesian; clicks were excluded by scope. Each language is audited under one declared reference variety (General American English; Standard Mandarin; Delhi Hindi; Castilian Spanish; Parisian French; MSA; Brazilian Portuguese; Seoul Korean; Hanoi Vietnamese; and the respective standards elsewhere), with inventories taken from the standard reference descriptions (for Slavic, Comrie & Corbett 1993). Every phoneme of each inventory is classified three ways: *exact* (a rendering that preserves all of that language’s contrasts involving the phoneme), *conventional merger* (a declared many-to-one mapping, enumerated below), or *declared loss* (prosodic folding). A rendering that depends on position (*/w/*, */ɲ/*) counts as exact when the positional rule is stated. The system is deterministic relative to this input: variety choice is a parameter of the audit, not of the mapping. Against existing practice the relevant baselines are the national transcription traditions — Bulgarian’s prescriptive system (Данчев 1982) and its analogues — which are recipient-adapting by design; Chuzhditsa differs precisely in preserving source contrasts those systems merge, which is what the citation-register framing predicts. The audit’s full 126-item example set, and the machine-readable phoneme-to-grapheme mapping it instantiates, accompany the repository. Table 4 summarizes the audit per language: the reference variety, the extension letters the language actually exercises, the conventional mergers engaged, and the declared losses. The pattern worth reading off is that no language needs more than a small subset of the inventory, that the mergers cluster on rhotics and sibilant places, and that declared loss is confined to prosody.

Language	Reference variety	Extensions exercised	Mergers engaged	Declared losses
English	General American	ʔ ɕ ʒ ʃ ʈ ɬ ɰ Ɂ	rhotic → ɹ	—
Mandarin	Standard (Beijing)	ɥ ɣ ʰ tones	/tʂ tɕ/ → ɕ; /z/ → ʃ	tone sandhi
Hindi–Urdu	Delhi	ʈ ʈʰ ɖ ɖʰ ɽ ɽʰ ɳ ɰ	/ɦ/ → ɦ	—
Spanish	Castilian	ɕ ʒ ʃ ʎ	/β/ → ɸ; rhotics	—
French	Parisian	ʌ ʌʱ ɔ ɔʰ ɔ̃ ɔ̃ʰ ɔ̃ʱ ɔ̃ʱʰ	/ʁ/ → ɹ	—
Arabic	MSA	ħ ʕ ʔ ʕ ʔ ʕ ʔ ʕ	/χ/ → ʁ	—
Bengali	Standard	ʈ ʈʰ ɬ	—	—
Portuguese	Brazilian	ɕ ʌ	/ʁ/ → ɹ	—
Russian	Standard	ы (native)	/ɕ:/ → ɕɕ	reduction (phonemic input)
Japanese	Tokyo	ɰ ʔ ɥ	/ɸ/ → ɸ; /w/ → ɰ	pitch accent
German	Standard	ö ʏ ɰ ɰ + xɐ	—	—
Korean	Seoul	ʰ ɰ ɰ ɰ	tense = geminate	—
Vietnamese	Hanoi	ɰ ɰ tones	/b d/ → ɸ ʌ	6 tones → 4 marks
Turkish	Istanbul	ö ʏ ɰ ɰ	/h/ → ʌ	—
Italian	Standard	s	—	—
Persian	Tehran	ɕ ɰ ɰ	/g/ → ɕ	—
Swahili	Standard	ɕ ʒ ʃ ɰ ɰ	implosives → ɸ ʌ	—
Greek	Standard Modern	ɕ ʒ ʃ	—	—
Polish	Standard	ʌ ʌʱ ʏ ɰ	/ɕ/ → ɕɕ; /ʂ/ → ɰ	—
Indonesian	Standard	ɰ ɰ ɰ	—	—

Table 4: Per-language audit summary. “Extensions exercised” lists the letters and marks the language’s inventory actually requires; palatalization via ɰ and length via ʰ are counted where load-bearing. Mergers and losses are the declared ones of §9.

Table 5 samples the audit; the full 126-item set accompanies the repository. Declared mergers are: all rhotics → **p**; /s//ʃ/ → **ш**; /f//f/ → **ф**; /h//h/ → **н**; implosives → **б а**; /q//g/ → **к**; Korean tense consonants as geminates; English input referenced to a rhotic accent (GA); /r/ is written only in Arabic material while /ʀ/ always is — a practical asymmetry in the laryngeal series, accepted rather than disguised; /r/ dropped outside Arabic material.

Source		IPA	Chuzhditsa
English	<i>thanks</i>	[θæŋks]	џАҢКС
English	<i>weekend</i>	[ˈwi:kend]	УЙКЕНА
Mandarin	北京 <i>Běijīng</i>	[peɪ.tɕiŋ], T3+T1	ПЭЙЧИН
Hindi	ठीक <i>thīk</i>	[tʰi:k]	ТʰИК
Arabic	محمد <i>Muhammad</i>	[muˈħammad]	МУХАММАД
French	<i>croissant</i>	[kʁwasɑ̃]	КРҰАСА
French	<i>bien</i>	[bjɛ̃]	БЯ
Polish	<i>Wrocław</i>	[ˈvrɔʦwaf]	ВРОЦҰАФ
Japanese	東京 <i>Tōkyō</i>	[to:kʰo:]	ТОКЬО
Korean	한글 <i>hangul</i>	[hangwɭ]	НАНГЫЛ
Georgian	წყალო <i>c'qali</i>	[tsʰqʰali]	ЦҚІАЛИ
Albanian	<i>Shqipëria</i>	[ʃcipəˈria]	ШКЫПЬРІА

Table 5: Sampled audit items; the Chuzhditsa column is set in the typeface under discussion, with live GSUB fusion and GPOS anchoring. In the Mandarin row the caron on Ě is tone 3, attached to the first syllable’s vowel; tone 1 is unmarked by rule.

10 Adoption: lessons from planned and organic systems

A designed register lives or dies socially, not technically; the sociolinguistics of script choice behaves like the sociolinguistics of language choice (Fishman 1977; Unseth 2005). Four historical cases bound the space Chuzhditsa enters. The comparison is organized by recurrent factors rather than narrative sympathy: community ownership, institutional sponsorship, literacy burden, identity payoff, compatibility with existing orthographic practice, domain of use, prestige and stigma, and the mundane infrastructure of input methods and type.

The planned failures. Volapük collapsed within a decade, less on linguistic defects than on its inventor’s proprietary control of the standard (Garvía 2015). Esperanto — maximally learnable, energetically organized — plateaued as a voluntary subculture: it extended nobody’s existing practice and attached to no prior community, so adopting it meant *joining a new identity* rather than exercising an old one (Forster 1982; Garvía 2015). The Deseret alphabet is the sharpest lesson, because it controls for institutional power: commissioned by Brigham Young, taught in classes from church-printed primers, backed by church funds, it was abandoned within roughly two decades — sponsorship without a felt need or an identity payoff bought nothing (Alder, Goodfellow & Watt 1984). Learnability and patronage, the two levers a designer can pull, decided none of these cases.

The designed success. Hangul is the one engineered script that won, and its path is instructive precisely because the design was not what won. Completed in 1443 and promulgated in 1446 (Ledyard 1998), it was opposed at once by the literati — the collective Ch’oe Malli memorial of 1444 — and spent four and a half centuries outside prestige writing, surviving in low-stakes registers: women’s correspondence, vernacular fiction, private devotion (Lee & Ramsey 2011). What revalued it was not its featural elegance but nineteenth-century nationalism, culminating in the 1894 Kabo edict that made the national script the base of official documents (King 1998). The mandate arrived 448 years after the design; the niche did the surviving.

The organic success. Katakana never asked anyone to change anything. It began as monks’ abbreviated annotation glosses on Chinese texts in the early Heian period, drifted into a division of labor

with hiragana over a millennium, and was codified only in 1900 by the implementing regulations of the Elementary School Order — with the loanword-register specialization consolidating later still (Seeley 1991; Twine 1991). Codification ratified practice; it did not create it.

The pattern reduces to four conditions, against which Chuzhditsa can be scored honestly. (i) *Additive, not substitutive*: the winners changed no existing spelling; Chuzhditsa alters nothing in any native orthography by construction. (ii) *Identity attachment*: Hangul survived on Korean-ness, not on design merit; Chuzhditsa’s revivals — the yuses, the Glagolitic letterforms — are the analogous move available to a designed system, heritage rather than invention. (iii) *A survivable low-stakes niche before any mandate*: Hangul had the vernacular genres, katakana the annotation margin; the dictionary pronunciation field, the language classroom, and the subtitle are proposed as the equivalent low-stakes niche. (iv) *Codification last*: every mandate that worked ratified existing practice, which is why this specification petitions no orthographic authority. Four cases cannot license induction, and none of this predicts uptake; the comparison identifies necessary-looking conditions, not sufficient ones. One condition, moreover, cannot be engineered around: on the planned/organic axis Chuzhditsa sits with Esperanto and Deseret, not with katakana. Assembly from attested material is a mitigation — every letter arrives with a history already attached — but it is not an exemption, and the Deseret case fixes an upper bound on what design quality plus sponsorship can buy without a community that recognizes itself in the script.

11 Limitations

Register versus encoding. The register’s marking lives at two levels that must not be conflated. The phonological extensions are plain-text-safe: they are standard Unicode letters that survive copy, paste, search, and font substitution. The *visual* register, however, is carried by typeface choice — a diaphasic diglyphia in Bunčić’s terms, not an encoded script contrast — and it degrades exactly where font control degrades: plain-text channels keep the letters but lose the look. Whether readers actually perceive the Glagolitic-derived forms as “foreign register” rather than “archaic,” “ecclesiastical,” or merely “a display font” is an empirical question this paper does not answer; a perception study — mixed native/Chuzhditsa text, register-identification and readability tasks with Cyrillic readers — is the single most valuable next experiment.

Input assumptions. The mapping consumes a source pronunciation in a declared reference variety. Dialect choice, phonemic versus phonetic input, and spelling-influenced variant pronunciations are parameters the user supplies, not problems the system solves; learnability of the extension letters, and the cost of input methods for them, are likewise open. The keyboard prototype in the repository addresses input mechanics but not habit.

System-internal residue. The retraction hook is polysemous by direction — uvular on **к**, pharyngeal on **х**, interdental on **ц** **з**, retroflex on **т** **а** **п**. No attested language contrasts two of its values on one base letter, so the ambiguity is theoretical, but it is a real weakening of the one-mark-one-meaning claim. **/с/** keeps two spellings (**сб** under contrast, **ш** by entrenchment) — the one point where usage beat system. Six-tone prosodies fold lossily into four marks. Click phonology is future work, as is a mark-to-mark (**mkmk**) positioning pass for the rare double-stacked vowel.

12 Conclusion

Chuzhditsa demonstrates that a katakana-like visual register, combined with deterministic phonological citation from a specified source pronunciation, can be retrofitted onto an alphabet without inventing a parallel inventory: a featural grammar over attested Cyrillic material covers twenty phonologies, one historical script supplies the register aesthetics, and the whole specification compiles to a font whose OpenType rules *are* the orthography. The generalization from Bulgarian to the Slavic superset cost

one letterform (e) and no grammar changes, which we take as evidence that the featural layer, not the letter list, is the right unit of design.

References

- [1] Alder, D. D., P. J. Goodfellow & R. G. Watt. 1984. Creating a new alphabet for Zion: The origin of the Deseret alphabet. *Utah Historical Quarterly* 52(3). 275–286.
- [2] Bunčić, D., S. L. Lippert & A. Rabus (eds.). 2016. *Biscriptality: A Sociolinguistic Typology*. Heidelberg: Universitätsverlag Winter.
- [3] Cahill, M. & K. Rice (eds.). 2014. *Developing Orthographies for Unwritten Languages*. Dallas, TX: SIL International.
- [4] Comrie, B. & G. G. Corbett (eds.). 1993. *The Slavonic Languages*. London: Routledge.
- [5] Comrie, B. 1996. Adaptations of the Cyrillic alphabet. In P. Daniels & W. Bright (eds.), *The World's Writing Systems*, 700–726. New York: OUP.
- [6] Coulmas, F. 2003. *Writing Systems: An Introduction to their Linguistic Analysis*. Cambridge: CUP.
- [7] Cubberley, P. 1996. The Slavic alphabets. In P. Daniels & W. Bright (eds.), *The World's Writing Systems*, 346–355. New York: OUP.
- [8] Данчев, А. 1982. *Българска транскрипция на английски имена*. 2. изд. София: Народна просвета.
- [9] Daniels, P. & W. Bright (eds.). 1996. *The World's Writing Systems*. New York: OUP.
- [10] Fishman, J. A. (ed.). 1977. *Advances in the Creation and Revision of Writing Systems*. The Hague: Mouton.
- [11] Forster, P. G. 1982. *The Esperanto Movement*. The Hague: Mouton.
- [12] Garvía, R. 2015. *Esperanto and Its Rivals: The Struggle for an International Language*. Philadelphia: University of Pennsylvania Press.
- [13] Irwin, M. 2011. *Loanwords in Japanese*. Amsterdam/Philadelphia: John Benjamins.
- [14] Kang, Y. 2011. Loanword phonology. In M. van Oostendorp, C. J. Ewen, E. Hume & K. Rice (eds.), *The Blackwell Companion to Phonology*, vol. 4, 2258–2282. Malden, MA: Wiley-Blackwell.
- [15] Kim-Renaud, Y.-K. (ed.). 1997. *The Korean Alphabet: Its History and Structure*. Honolulu: University of Hawai'i Press.
- [16] King, R. 1998. Nationalism and language reform in Korea: The questione della lingua in precolonial Korea. In H. I. Pai & T. R. Tangherlini (eds.), *Nationalism and the Construction of Korean Identity*, 33–72. Berkeley: Institute of East Asian Studies.
- [17] Ledyard, G. K. 1998. *The Korean Language Reform of 1446*. Seoul: Sin'gu Munhwasa.
- [18] Lee, K.-M. & S. R. Ramsey. 2011. *A History of the Korean Language*. Cambridge: CUP.
- [19] Peperkamp, S. & E. Dupoux. 2003. Reinterpreting loanword adaptations: The role of perception. *Proceedings of the 15th International Congress of Phonetic Sciences*, 367–370.

- [20] Robertson, W. C. 2020. *Scripting Japan: Orthography, Variation, and the Creation of Meaning in Written Japanese*. London: Routledge.
- [21] Sampson, G. 1985. *Writing Systems: A Linguistic Introduction*. Stanford: Stanford University Press.
- [22] Schenker, A. M. 1995. *The Dawn of Slavic: An Introduction to Slavic Philology*. New Haven: Yale University Press.
- [23] Sebba, M. 2007. *Spelling and Society: The Culture and Politics of Orthography around the World*. Cambridge: CUP.
- [24] Seeley, C. 1991. *A History of Writing in Japan*. Leiden: Brill.
- [25] Sproat, R. 2000. *A Computational Theory of Writing Systems*. Cambridge: CUP.
- [26] Stanlaw, J. 2004. *Japanese English: Language and Culture Contact*. Hong Kong: Hong Kong University Press.
- [27] Twine, N. 1991. *Language and the Modern State: The Reform of Written Japanese*. London: Routledge.
- [28] The Unicode Consortium. 2024. *The Unicode Standard*, Version 16.0.0. South San Francisco, CA: The Unicode Consortium.
- [29] Unseth, P. 2005. Sociolinguistic parallels between choosing scripts and languages. *Written Language & Literacy* 8(1). 19–42.