

Chuzhditsa: a featural loanword register for Cyrillic, from Bulgarian to the Slavic superset

Author name withheld for review*

Abstract

Japanese distinguishes loanwords at a glance: katakana is a parallel register whose visual form alone announces foreignness, and whose adaptation rules map any foreign word to exactly one spelling. No Cyrillic orthography has an equivalent, and the result is systematic indeterminacy (Bulgarian УОШИНГТОН ~ ВОШИНГТОН; Serbian and Russian counterparts abound). We present *Chuzhditsa* (Bulg. чуждица, ‘foreignism’), a three-layer extension of Cyrillic that supplies (i) a featural grammar of diacritic operations deriving letters for non-Slavic phonemes, (ii) a prosodic layer for stress, tone, and length, and (iii) a register-marking typeface whose Glagolitic-derived forms perform the katakana function visually. The letter inventory is assembled conservatively: revivals from historical Slavic orthography (the yuses А а, dzelo С с), imports from the wider Cyrillic ecosystem (Kazakh Қ қ, Bashkir Җ җ, Tajik Ў ў), and rule-derived residue. We evaluate the system against the phonologies of twenty major languages, generalize it from its Bulgarian pilot to the full Slavic Cyrillic superset (with deterministic bridges from Latin-alphabet Slavic languages and an external validation on Albanian), and describe a parametric implementation as an OpenType typeface family in which the fusion and anchoring rules of the orthography are executable GSUB/GPOS code. We argue that the laryngeal series (Х х Һ plus the aspiration modifier Ҫ ҫ) illustrates a general principle for extending alphabets: articulatory gradients should map to visual gradients.

1 Introduction

Every literate community borrows words faster than its orthography can naturalize them. Japanese is unusual in having institutionalized the problem: katakana is a full parallel syllabary whose use *is* the marking of foreignness, and whose adaptation conventions ([nart] → ナイト) are deterministic enough that two writers converge on one spelling. The Cyrillic-writing world has no analogous device. Loanwords enter Bulgarian, Russian, Serbian or Ukrainian through ad hoc transcription: sound-based and spelling-based renderings compete (ЛАНЧ ~ ЛЪНЧ for *lunch*), phonemes without Cyrillic letters collapse unpredictably (/θ/ → Т or С; /w/ → В or Y), and nothing in the written form signals that a word is a guest rather than a native.

This paper describes Chuzhditsa, a designed orthographic register that supplies the missing device for Cyrillic, piloted on Bulgarian and then generalized to the Slavic superset. The system makes three claims of general interest to writing-system design. First, that a loanword register can be *featural* in Sampson’s (1985) sense: rather than inventing letters ad libitum, a small set of diacritic operations — each meaning exactly one phonetic feature — derives the entire extension inventory, echoing the compositional logic of Hangul. Second, that such a register can be assembled almost entirely from attested material: dead letters of the target orthography’s own history and living letters of related orthographies, with invention reserved for genuine residue. Third, that the register-marking function that Japanese distributes over a second character inventory can instead be carried by a second *typeface register* — realized here as an OpenType family whose substitution (GSUB) and positioning (GPOS) rules are the orthography’s fusion and attachment rules, executable rather than merely described.

*Fonts, build toolchain, and the twenty-language audit are available in the accompanying repository; the typeface used for all object-language examples in this manuscript is the system being described.

2 Background

2.1 Katakana as an adaptation layer

Katakana’s success rests on three properties (Irwin 2011): a fixed sound-mapping table applied to the source word’s pronunciation, not its spelling (silent letters do not survive); graphic distinctness sufficient to mark register at a glance; and extensibility (ㇰ for /v/, ㇱ ㇲ for /fa/) when the inventory runs out. Chuzhditsa adopts all three as requirements.

2.2 The Cyrillic commonwealth

Cyrillic serves some fifty languages, and the non-Slavic ones have already engineered letters for most sounds Slavic lacks: Kazakh ҕ /q/, Һ /h/, Җ /ɣ/; Bashkir Ҙ /θ/, ҙ /ð/; Tajik Ҳ /ħ/, Ӣ /i:/; Altai ӧ /ø/, ӧ̈ /y/; Belarusian ӱ /w/; Serbian џ /dʒ/; the Caucasian palochka I. These letters carry decades of typographic and pedagogic practice. A design that imports them inherits that practice for free — an option unavailable to katakana’s designers and, we will argue, the single largest economy in the present system.

2.3 The Glagolitic precedent

Bulgaria is the locus of a script replacement: Glagolitic (863 CE), the first Slavic alphabet, was superseded at Preslav (c. 893) by Cyrillic, which retained Glagolitic-derived shapes only for sounds Greek lacked (Schenker 1995; Cubberley 1996). Glagolitic then survived for a millennium as a special-register script (Croatian liturgical use into the nineteenth century). Chuzhditsa knowingly reverses the vector: the visual DNA of the *oldest* Slavic letterforms — round bowls, unicameral proportions — marks the *newest* words, and the letters it revives (А а С с Ѧ ѧ) all have Glagolitic ancestors.

3 Design principles

1. **Determinism.** One phoneme, one rendering; two transcribers converge. Loss is permitted but must be declared (§9).
2. **Featural transparency.** Every diacritic operation means exactly one feature and applies wherever the feature appears.
3. **Pair symmetry.** Obstruent letters enter the system in voiceless/voiced pairs aligned with the native series. This principle outranks heritage: the historical theta-letter fita (Θ) was rejected because it has no voiced partner and derives from no rule; the Bashkir pair Ҙ/ҙ (/θ//ð/) does its work systematically.
4. **Precedent first.** Revive native dead letters, then import living Cyrillic letters, and only then derive new combinations; never invent an unattested base shape.
5. **Register by typeface.** Foreignness is marked by letterform register, not by punctuation or a second inventory: text set in the Chuzhditsa face is thereby marked, exactly as katakana text is.

4 The featural grammar

Table 1 gives the full operation set. The retraction hook generalizes a precedent internal to the Slavic base inventory itself — Ң is Ң with a hook — and the glide breve generalizes Ѧ → ѧ. Where Unicode supplies precomposed hooked letters (Ҙ Һ Җ Ҙ Җ) they are used; elsewhere the same hook is written with combining U+0322 (Ṛ Ṛ p).

drops length first (Ѣ). And four marks cannot represent six-tone systems without folding (Vietnamese *hỏi/ngã* merge). The design verdict generalizes katakana’s own: prosody is a dictionary layer, not an orthographic one.

7 From Bulgarian to the Slavic superset

The pilot treated Bulgarian’s thirty letters as the base. Nothing in the featural grammar is Bulgarian-specific, however, and the generalization is administrative rather than structural: the base becomes each language’s own Cyrillic alphabet, and the extension and prosodic strata apply unchanged. Concretely, the typeface’s character set was widened to the Slavic Cyrillic superset — Russian Ѣ Ѥ Ѧ, Ukrainian і ї є ґ, Belarusian ґ, Serbian ј ѡ Ѣ ѣ Ѧ ѧ, Macedonian ѕ ґ ґ — so that any Slavic Cyrillic text sets natively. One design conflict surfaced and is instructive: the pilot’s round e-form was, in fact, the shape of Ukrainian є; pan-Slavic coverage forced e to an angular form so that the e/е contrast survives. Extending a script for one language can silently occupy a sibling’s letterform space; superset-first design catches this early.

For each language the system’s marginal contribution differs, and several languages turn out to need the extension layer for *native* material, not only loans: Slovak *ä* is Ä and Slovak *h* is /h/ → ѧ; Slovenian’s native schwa (*polglasnik*) is exactly Ѣ (пѣс /pəs/ ‘dog’); Serbian’s pitch accents are the prosodic layer’s use case; Polish *q* ę are the yuses meeting their own descendants (макѣ, пѣчѣ). Latin-alphabet Slavic languages connect through a deterministic bridge: Gaj’s digraphs map one-to-one onto Serbian Cyrillic letters (*lj nj dž* → ѡ Ѣ ѧ), and the residue is featural — Czech ř /r/, the notorious singleton, is written ř with the caron it already wears; Slovak long syllabic liquids *ĺ ř* take the length macron (ѡѣѣ /vr:ba/).

As an external stress test we applied the system to Albanian, a non-Slavic Balkan language: every Albanian phoneme lands on an existing letter (ë /ə/ → Ѣ, th/dh → ѣ/ѧ, y → Ѥ, xh → ѧ, x /dz/ → ѣ, q/gj → ѣ/ѧ): Шкѣпѣрѣѣ, Тѣрѣѣѣ. A system built for Slavic phonology plus the featural operations appears to cover the Balkan sprachbund without amendment.

8 Typeface implementation

The Chuzhditsa family (Regular, Bold, Italic, Bold Italic; TTF and CFF/OTF) is generated by a ~300-line parametric builder: every glyph is a set of monoline strokes — lines, arcs, and the full circles that give the face its Glagolitic character — expanded to outlines by a union of per-segment rectangles and per-vertex discs, with weight and slant as parameters. Capitals stand at 700 units, lowercase at an x-height of 500 with true two-case architecture (bowl-and-stem ѣ, ascending ѧ, descending ѧ). Every combining mark receives computed GPOS mark-to-base anchors (the macron of ѧ attaches 251 units lower than that of Ä; the retroflex hook of ѧ lands on the letter’s stem), and the ligature stratum is a GSUB *liga* feature — the orthography’s rules are the font’s code, verified by HarfBuzz shaping tests in continuous integration. The manuscript the reader is holding sets all object-language material in this typeface.

9 Evaluation: the twenty-language audit

The system was audited against the phoneme inventories of the twenty most-spoken languages (clicks excluded by scope). Table 2 samples the audit; the full 126-item set accompanies the repository. Declared mergers are: all rhotics → ѧ; /s//ʃ/ → ѧ; /f//f/ → ѧ; /h//h/ → ѧ; implosives → ѧ ѧ; Korean tense consonants as geminates; /r/ dropped outside Arabic material.

Source		IPA	Chuzhditsa
English	<i>thanks</i>	[θæŋks]	ЇАҢКС
English	<i>weekend</i>	[ˈwi:kɛnd]	ЎЙКЕНА
Mandarin	北京 <i>Běijīng</i>	[pèi.tɕiŋ]	БЭЙЦИҢ
Hindi	ठीक <i>thīk</i>	[tʰi:k]	Т'ЙК
Arabic	ﻣﯘﮬﻤﻤﺎﺩ <i>Muḥammad</i>	[muˈħammad]	МУХАММАД
French	<i>croissant</i>	[kʁwasɑ̃]	КРЎАСА
French	<i>bien</i>	[bjɛ̃]	БЮ
Polish	<i>Wrocław</i>	[ˈvrɔʦwaf]	ВРОЦЎАФ
Japanese	東京 <i>Tōkyō</i>	[to:kʰo:]	ТОКЬО
Korean	한글 <i>hangul</i>	[hangwɭ]	НАНГЫЛ
Georgian	ცყალი <i>c'qali</i>	[tsˈqʰali]	ЦҚАЛИ
Albanian	<i>Shqipëria</i>	[ʃcipɐˈria]	ШКЫПЕРИА

Table 2: Sampled audit items; the Chuzhditsa column is set in the typeface under discussion, with live GSUB fusion and GPOS anchoring.

10 Limitations

The retraction hook is polysemous by direction — uvular on **к**, pharyngeal on **х**, interdental on **с** **з**, retroflex on **т** **д** **р**. No attested language contrasts two of its values on one base letter, so the ambiguity is theoretical, but it is a real weakening of the one-mark-one-meaning claim. /с/ keeps two spellings (**сѣ** under contrast, **ш** by entrenchment) — the one point where usage beat system. Six-tone prosodies fold lossily into four marks. Click phonology is future work, as is a mark-to-mark (mkmk) positioning pass for the rare double-stacked vowel.

11 Conclusion

Chuzhditsa demonstrates that the katakana function — deterministic adaptation plus visible register — can be retrofitted onto an alphabet without inventing a parallel inventory: a featural grammar over attested Cyrillic material covers twenty phonologies, one historical script supplies the register aesthetics, and the whole specification compiles to a font whose OpenType rules *are* the orthography. The generalization from Bulgarian to the Slavic superset cost one letterform (e) and no grammar changes, which we take as evidence that the featural layer, not the letter list, is the right unit of design.

References

- [1] Comrie, B. & G. Corbett (eds.). 1993. *The Slavonic Languages*. London: Routledge.
- [2] Coulmas, F. 2003. *Writing Systems: An Introduction to their Linguistic Analysis*. Cambridge: CUP.
- [3] Cubberley, P. 1996. The Slavic alphabets. In P. Daniels & W. Bright (eds.), *The World's Writing Systems*, 346–355. Oxford: OUP.
- [4] Данчев, А. и колектив. 1989. *Българска транскрипция на английски имена*. София: Народна просвета.
- [5] Daniels, P. & W. Bright (eds.). 1996. *The World's Writing Systems*. Oxford: OUP.
- [6] Irwin, M. 2011. *Loanwords in Japanese*. Amsterdam: John Benjamins.
- [7] Sampson, G. 1985. *Writing Systems*. Stanford: Stanford University Press.

- [8] Schenker, A. 1995. *The Dawn of Slavic*. New Haven: Yale University Press.
- [9] Sproat, R. 2000. *A Computational Theory of Writing Systems*. Cambridge: CUP.
- [10] The Unicode Consortium. 2024. *The Unicode Standard*, Version 16.0. Mountain View, CA.